Queries that work well:
"ics 32 python projects requirements"
"bren hall map"
"cs 122a project"
"ACM"
"ics capstone examples"
"graduate admissions"
"cristina lopes"
"women in cs"
"hello"
"python courses"

Queries that return duplicates:
"cs advisor appointment"
"artificial intelligence"
"ai"
"graduate admissions"
"office of uci ics"

Queries that are too slow:
"master of software engineering"
"uci ics computer research"
"computer programming data"
"research in graduate cs topics"
"ai cs professor phd researcher"

To fix the problem of duplicates, I decided to add exact duplicate detection. To fix the problem of slow queries I had to change the way I store the tokens in partial indices, as the first few were too big and slowed the code down too significantly.